

# Building an Investment Approval Workflow Application on Replit – Implementation Guide

## Functional Requirements Overview (Tawuniya RFP Item 1)

According to Tawuniya's RFP Item 1, the investment team needs an automated workflow system for two primary processes: **(a)** creating a new investment opportunity, and **(b)** requesting cash for an investment <sup>1</sup> <sup>2</sup>. In addition, the system must support **SLA tracking**, **document uploads**, and **multi-tier approvals** to ensure that all necessary approval stages are covered <sup>3</sup> <sup>4</sup>. Below we break down these functional requirements:

### New Investment Creation Workflow

- **Purpose:** Enable employees (investment analysts) to create a new investment request (opportunity) with all required details and route it through the necessary approval stages <sup>3</sup>. This automates the initiation of a new investment proposal.
- **Key Steps:**
  - **Create Request:** The user fills out a "New Investment" form with details such as the target investment company, investment type (asset class), proposed value/amount, identified risks, and any other required data <sup>3</sup>. The user should also attach all relevant documents (e.g. investment proposal PDF, financial analyses) needed for approvals.
  - **Submission & Routing:** Upon submission, the system records the request and triggers the **approval workflow**. The request is assigned a unique ID and initial status (e.g. "Pending Approval"). It generates tasks for the first approver in the chain.
  - **Multi-Stage Approval:** The request passes through **tiered approval levels** ("all stages of approvals in the company" <sup>3</sup>). For example, first a line manager approval, then a department head or investment committee, etc., depending on configured rules. Each approver reviews the investment details and attached documents and can approve or reject.
  - **Tracking & Updates:** The requestor can **follow up on the status** of their submitted request at any time <sup>5</sup>. They should see the current approval stage, who it's awaiting action from, and the *expected completion time* for the next step (if defined by SLA).
  - **Editing & Resubmission:** If an approver requests changes or more info (or rejects with comments), the workflow might allow the requestor to **edit the submitted request** and resubmit it <sup>5</sup>. The system should maintain version history or comments for audit trail.
- **Completion:** Once all required approval tiers have approved, the investment request is marked **Approved** and becomes an active investment record. A final notification is sent to the requestor and relevant stakeholders.

- **SLA Enforcement:** The system should track the time each approval stage takes against defined **Service Level Agreements (SLAs)** <sup>4</sup> . If a request remains “pending” beyond the agreed SLA for a stage, automatic reminders/notifications are sent to the responsible approver and possibly escalated to supervisors. For example, if an approval is expected within 2 business days and it’s overdue, the system notifies the approver and CC’s an admin.
- **Pending Requests Visibility:** Users need to easily see a list of their “open/pending” investment requests. As per the RFP, employees can view “**pending open requests and search past requests**”, and if a request is not closed within the SLA, the employee receives a **notification** reminding them (or the responsible party) to close it <sup>4</sup> .
- **Templates:** To streamline frequent submissions, the system should allow **request templates**. This means users can load a predefined form layout or defaults for common investment types, reducing data entry for repetitive fields <sup>4</sup> .

## Cash Request Workflow

- **Purpose:** Enable an employee to request a cash disbursement or funding for an already-approved investment, which then goes through approval and notifies Finance <sup>2</sup> . For instance, after an investment is approved, the investment team might need to draw cash to fund that investment.
- **Key Steps:**
  - **Create Cash Request:** The user selects the relevant **investment** (from a list of active investments) and enters details like the cash amount needed, reason/purpose of the cash request, and any supporting documents (e.g. invoice, payment schedule).
  - **Submission & Approval:** The cash request is routed to the appropriate **manager for approval** <sup>2</sup> . Typically, this would be a manager in the investment department or the owner of the investment budget. The manager reviews the request details and either approves or rejects.
  - **Finance Notification:** Once approved by the manager, the system sends a **notification to the Finance department** (or a designated Finance user) <sup>2</sup> . This could create a task for Finance to process the payment. (Finance may not need to “approve” in the system, unless a secondary approval is required; the RFP implies Finance is just notified.)
  - **Tracking & Completion:** The requester can track the status of the cash request similar to investment requests. After manager approval (and possibly Finance processing), the cash request is marked completed (funds released).
- **Similar Features:** The cash request workflow shares many of the same supporting features as the investment workflow:
  - **Follow-up and Edit:** The requester can follow the status and edit the request if needed (e.g., if the manager requests changes) <sup>6</sup> .
  - **Task List:** Approvers see the cash request in their task list for action <sup>7</sup> .
  - **SLA Tracking:** If cash approvals have SLA (e.g., must be approved within X hours/days), the system should send reminders for delays <sup>8</sup> (just as with investment requests).

- **Templates:** Possibly templates for recurring cash requests (less common, but the system supports template creation as well) <sup>8</sup>.

## Investment Management Page (Post-Approval)

Once an investment request is approved and becomes an active investment, the system should provide an **Investment Management page** for ongoing management <sup>9</sup>. This page is essentially a detailed view of the investment and related activities:

- **Investment Details & Documents:** It shows all details of the investment (from the creation form) and all documents related to it (proposal documents, contracts, etc.) <sup>10</sup>.
- **History Timeline:** The page should display a **history or timeline** of all actions and updates on this investment <sup>11</sup> – e.g., submission date, approval dates by each approver, any modifications, and cash requests linked to it.
- **Post-Approval Actions:** Users with the right permissions can manage the investment here – e.g., update certain fields (if allowed), upload additional documents (like reports or updates), or even initiate sub-workflows like **cash requests** directly from this page <sup>12</sup>.
- **Cash Requests Listing:** All **cash requests related to this investment** should be visible here for context <sup>12</sup>. For example, a section of the page could list each cash drawdown request, its status, and amount.
- **Permissions:** Typically, only certain roles (investment managers or admins) can modify an active investment record, whereas others may have read-only access. The page should enforce those permissions.

## Task Management and Notifications

A core requirement is to have robust **task management** so that team members can easily see and act on their pending work <sup>5</sup> <sup>7</sup>:

- **My Tasks List:** Each user (especially managers/approvers) should have a “My Tasks” view showing all requests waiting for their action. This includes pending investment approvals, cash request approvals, or any other workflow tasks (like providing additional info). From this list, the user can click a task to view the request details and then take action (approve/reject or other specified actions) <sup>13</sup>.
- **Delegation and Workload Balancing:** The RFP mentions “*automated request distribution between the team based on availability*” <sup>14</sup>. This implies that tasks can be auto-assigned to team members in a round-robin or load-balanced fashion. In implementation, this could mean:
  - If multiple people can handle a task (e.g. there are multiple analysts or approvers at a certain level), the system should assign it to an available person (perhaps the next in line or someone not overloaded).
  - An admin interface can allow managing these assignment rules or delegating tasks if someone is out of office.
- **Alerts and Notifications:** The system should send **notifications** for key events:
  - When a task is assigned to a user (e.g., “You have a new approval request to review”).
  - SLA breach alerts (e.g., “Investment Request #123 is overdue for approval”).
  - Completion notifications (e.g., “Your investment request has been approved” or “Cash request released”).

- These notifications could be emails or in-app notifications (or both). For in-app, a notification bell icon can show unread notifications. All notification events should be logged (for audit and for users to review past notifications).
- **SLA Monitoring Dashboard:** For transparency and compliance, an SLA dashboard can list all open tasks with their deadlines and highlight those approaching or exceeding SLA. This helps managers intervene proactively. This might be part of an admin or manager view.

## Document Upload and Management

Document handling is critical since each investment request may require multiple attachments. Key points:

- **Upload Feature:** Users must be able to upload supporting documents when creating a request (e.g., investment proposals, risk assessments) or later attach documents to an existing request/investment (e.g., updated documents or cash request receipts) <sup>3</sup> <sup>15</sup>. The UI should allow multiple files to be attached, with a clear indication of allowed file types and size limits (to prevent overly large files).
- **Storage:** Rather than storing files in the Replit container (which is ephemeral and limited), files should be stored in a persistent **Object Storage**. Replit provides a built-in Object Storage feature backed by Google Cloud Storage <sup>16</sup>. Each file is stored as an object in a bucket, ensuring high durability and availability. The system will keep only metadata in the database (like a file URL or key, file name, uploader, timestamp, etc.).
- **Access Control:** Documents might contain sensitive data. Implement proper access controls:
  - Only authorized users can download/view the documents. For instance, only approvers or involved parties can access confidential proposal documents.
  - Generate signed URLs or use a secure proxy route for file download to ensure only authenticated users with permissions get the file.
- **Versioning:** Consider storing version history if users upload updated versions of documents (or at least keep old files rather than overwriting, for audit purposes).
- **Virus Scanning:** For security, it's a good practice to scan uploaded files for viruses or malware. This could be an integration point (perhaps using a third-party scanning API) if required by corporate security policies.
- **Metadata:** The UI should display attached documents with metadata like name, type, size, and an option to download. Optionally, allow users to add a description for each file (e.g., "Financial Model" or "Risk Analysis Report").

## Tiered Approval Process (Multi-Level Workflows)

The application must support **tiered approvals**, meaning a request can require multiple sequential approvers. As noted, a new investment may need to pass “**all stages of approvals in the company**” before it's fully approved <sup>3</sup>. Implementation considerations:

- **Defining Approval Chains:** For each request type (Investment or Cash), define what the approval chain looks like. This could be configured via an admin interface or configuration file. For example:
  - Investment Request: Level1 = Investment Manager, Level2 = Investment Committee, Level3 = COO/ CFO (if above certain amount).
  - Cash Request: Level1 = Department Manager, Level2 = Finance Dept (for final processing).
  - The workflow rules might depend on attributes (e.g., if amount > X, add an extra approval from Risk or CFO).

- **Dynamic vs. Static Workflow:** Ideally the workflow should be **configurable** (the RFP suggests ability to change workflows via visual tools and rule-based configurations <sup>17</sup> ). For an MVP, you might hardcode or define in a config the required steps. Long term, storing these in a database (e.g., a `WorkflowDefinition` with steps) allows admins to modify without code changes.
- **Task Creation per Level:** When a request is submitted, the system creates an **approval task** for the first level approver. When that approver approves, the system automatically creates a task for the next level, and so on, until all levels completed. If any approver rejects, the workflow could either terminate (request marked Rejected) or pause and return to requester for changes (depending on business rules).
- **Parallel Approvals:** If any steps can be parallel (not mentioned in RFP but could be a need, e.g., two approvals needed but order doesn't matter), the system should handle creating multiple tasks at once for that level and requiring all to complete.
- **Escalation:** If an approver is unavailable or the task is overdue, there should be a way to **reassign or escalate** to an alternate approver. Perhaps the system automatically escalates to the approver's manager or a backup person after SLA breach.
- **Audit Trail:** Every approval decision (approve/reject) should be logged with who made the decision, timestamp, and any comments. This can be stored in an **Approvals** table (see schema below) separate from the main request, to maintain a history of all approval stages.

## User Roles and Permissions

The application will be used by different types of users, each with specific permissions (this corresponds to RFP's *"Types of users and their permissions"* under User Management <sup>18</sup> ). Proposed roles:

- **Investment Team Member (Analyst):** Can create new investment requests and cash requests, view their own requests, edit their requests (until final approval), view task assignments if they need to provide more info. They likely cannot approve requests.
- **Investment Manager:** Can create requests as well, but primarily is an approver for Investment Requests (first approval level). Can view requests from their team, and has a task list of pending approvals. May also initiate cash requests or approve cash requests, depending on company policy.
- **Department Head / Investment Committee:** Higher-level approvers for large investments. They receive tasks for second-level approvals. Their interface is similar to manager: a task list and ability to view and approve/reject.
- **Finance Officer:** Approver/processor for cash requests. They receive notifications when a cash request is approved by the manager. Possibly they mark the request as "Processed" once funds are transferred. They might also have read-only access to see investment details for context.
- **Administrator:** Can manage system settings, user accounts, roles, and workflow configurations. Admins can see all requests, reassign tasks, override in emergencies, and configure master data (like lookup values, SLA durations, etc.). They also manage templates and have access to audit logs.

*Permissions:* The system should enforce role-based access control: - Only approvers can see the "Approve/Reject" actions for tasks assigned to them. - Requestors can edit their requests only until submission (or until an approval is granted; after that, edits may need to create a change request). - Only Admins can access user management screens or modify workflow definitions. - Sensitive fields or actions should be visible/available only to appropriate roles (for example, maybe only Finance can see certain financial account info, etc., if applicable).

The Replit platform itself supports secure authentication and can integrate with single sign-on if needed (Replit offers SSO and RBAC features on enterprise plans) <sup>19</sup>. However, within our application, we will implement our own user accounts and roles, or integrate with the company's SSO if required. It's recommended to use standard authentication best practices (hashed passwords or OAuth2 flows) rather than managing plain passwords.

## Recommended Architecture on Replit (Frontend & Backend)

Building this as a **full-stack web application** on Replit is feasible. Replit provides an environment where you can run both the frontend (client-side code) and backend (server + database) within a single Replit workspace or as separate services. The architecture should be modular, separating concerns of UI, API, and database.

### Frontend Architecture

- **Technology Choice:** A modern **single-page application (SPA)** using a framework like **React** is recommended for a rich interactive UI. React (with libraries like Vite or Create React App) can be easily set up on Replit (there are templates for full-stack React+Node) <sup>20</sup>. Alternatively, other frameworks like Vue or Angular could be used, but React has a large ecosystem for enterprise UI components (Material-UI, Ant Design, etc.) which can accelerate development.
- **Structure:** Organize the React app with a routing system for different pages (e.g. using React Router). Use component-based design for forms and tables (e.g. a `<InvestmentForm />`, `<RequestList />`, `<TaskList />` etc.).
- **State Management:** Use a state management library (like Redux or React Context) to handle global state, such as the logged-in user info, and to temporarily store form data or lists of requests. This helps when multiple components need to share data.
- **UI Components:** Leverage a UI library for consistency and speed. For example, using **Ant Design** or **Material UI** can provide ready-made components for forms, tables, modals, notifications, etc. This will cover things like date pickers for SLA dates, file upload controls, and responsive layouts.
- **Design & Responsiveness:** Ensure the interface is intuitive for the end-user:
- Use a **dashboard layout** with a sidebar or menu for navigation (links to "My Requests", "My Tasks", "New Request", "Admin", etc., depending on role).
- Use clear forms with validation messages. For instance, if required fields are missing or numbers are out of range, show immediate feedback.
- Make it responsive so it's usable on various screen sizes (investment team might primarily use desktops, but some might check statuses on tablets).
- **Navigation Flow:** The frontend will manage client-side routing such that:
- After login, users land on a **Dashboard** (maybe showing a summary: count of pending tasks, recent requests, etc.).
- They can navigate to **Create Investment** or **Create Cash Request** forms.
- They can open **My Tasks** to approve requests assigned to them.
- They can open **My Requests** to see what they have submitted (with statuses).
- Admin users will have additional menu items like **Admin Panel** for user & workflow management.
- **API Integration:** The React app will communicate with the backend via **REST API** calls (or GraphQL if we prefer). For example:
- Submit form data to endpoints like `POST /api/investments`.

- Fetch lists with `GET /api/investments?filter=mypending` or `GET /api/tasks?assignee=me`.
- Use secure HTTP calls with the logged-in user's token for authentication.

## Backend Architecture

- **Technology Choice:** A robust backend framework is needed to handle business logic, workflows, and database interactions. **Node.js with Express** is a good option on Replit, especially since it pairs well with a React front-end and Replit templates support it. Alternatively, **Python with Flask or FastAPI** could be used if the team prefers Python – Replit can host either. For this guide, we'll assume a Node/Express stack (TypeScript would be beneficial for type safety in a project of this scope).
- **Layers:** Structure the backend in layers:
- **API Routes / Controllers:** Define Express routes for each entity (investments, cashRequests, tasks, etc.). These controllers handle HTTP requests, validate input, and call service logic.
- **Service Layer:** Implement core logic in services or use cases (e.g., an `InvestmentService` that contains methods like `createInvestment`, `approveInvestment`, etc.). This layer enforces business rules (like "if manager approves and amount > X, create an additional approval task for CFO").
- **Data Access Layer (DAL):** Use an ORM or query builder to interact with PostgreSQL. Replit's integrated database is PostgreSQL 15 (hosted on Neon) <sup>21</sup>. An excellent choice is **Prisma** or **TypeORM** (for Node) or **SQLAlchemy** (for Python). The ORM can map tables to classes, making it easier to manage relationships (like an `Investment` has many `Approval` records).
- **Workflow Engine:** To handle the multi-stage approvals, the backend can implement a simple workflow engine:
  - Define the approval chain (maybe as config or DB records).
  - When a request is created, instantiate a workflow: create initial tasks, etc.
  - A state machine pattern can be useful: e.g., an `InvestmentRequest` record has a status field that transitions from "Draft" → "Pending L1 Approval" → "Pending L2 Approval" → "Approved" (or "Rejected"). The service layer ensures that when one level approves, the status moves forward and next task is created.
  - Alternatively, leverage an existing workflow library if available. But custom logic is straightforward for two defined processes.
- **Notifications & Emails:** The backend should have a component for sending notifications. This could be as simple as sending emails using an SMTP service or using a service like Twilio SendGrid. For in-app notifications, you might have an `/api/notifications` endpoint and use WebSockets or periodic polling to let the front-end fetch new notifications for the user.
- **External Integrations:** If needed, integrate with company systems:
  - Single Sign-On (if employees should log in with corporate credentials).
  - Email server for notifications.
  - Logging/monitoring services for error tracking.
- But since the app is on Replit, you might keep it self-contained for initial implementation.
- **Scalability Consideration:** Node/Express can handle concurrent requests, but ensure to avoid blocking operations. Use async calls for DB and I/O. If heavy processing is needed (like generating PDF reports), consider offloading to background jobs or using Node worker threads. Replit's

environment could run background jobs as separate processes or threads if needed, but for an MVP, synchronous request/response with quick operations is fine.

## Database Integration (PostgreSQL on Replit)

Using a PostgreSQL database is recommended for relational data integrity and the complexity of entities we have. Replit makes it easy to add a PostgreSQL database to your app – it provides a **fully managed Postgres (via Neon)** which can be set up with one click or via the AI agent <sup>22</sup> <sup>23</sup>. Key points:

- **Schema Design:** (Detailed schema provided in a later section.) We will have tables for Users, Roles, Investments, CashRequests, Approvals, Tasks, Documents, Notifications, AuditLogs, etc. Ensure to define proper foreign keys (e.g., tasks reference the request they belong to, documents reference their parent entity, etc.) and use indices on frequently queried fields (like task assignee, request status).
- **ORM Usage:** If using an ORM like Prisma, define the data models there, and it will handle migrations. Otherwise, you can write SQL migrations to create tables. With Replit's DB tool, you can run SQL commands in the workspace to set up schema and also visually inspect data (Replit's database tool includes a visualizer for tables).
- **Connections:** The database connection string is provided via environment variables (`DATABASE_URL`). Replit securely stores these in Secrets so you don't expose credentials <sup>24</sup>. The backend should read this URL and connect using a standard Postgres client/ORM.
- **Performance:** For an enterprise workflow, ensure to use transactions where needed (e.g., when creating an investment request and initial tasks together) to maintain consistency. Also, consider query performance (e.g., loading a dashboard might require joining users, tasks, requests – those queries should be optimized with indexes).
- **Persistence:** The PostgreSQL instance is persistent (data won't be lost between restarts). This is crucial because Replit's filesystem is not persistent for long-term storage, so DB is the source of truth for data.

## Deployment Architecture on Replit

On Replit, you have a few options for running the app: - During development, you run the Express server (listening on some port) and the React dev server. In a single Repl, you can run both: e.g., run the React build script and serve static files via Express, or use two repls (one for front, one for back) during dev. - For simplicity, a common approach is to serve the frontend from the backend: - Build the React app into static files (HTML, JS, CSS) and have Express serve those from a `build` directory. - Thus you have one unified server (Express) that serves the React SPA and also serves the API endpoints. This avoids dealing with CORS and multiple servers. - Replit's default Node templates often do this for fullstack projects. - When ready to deploy (make the app continuously available to users): - Use **Always On** or **Reserved VM** on Replit to keep the server running. Replit free repls "sleep" after a period of inactivity, which is not acceptable for a production app. An Always On Repl (available with a paid plan or cycle boosts) will keep it running 24/7 <sup>25</sup> <sup>26</sup>. - For a more robust setup, Replit offers **Reserved VM Deployments**, which run your app on a dedicated always-on VM with reserved resources <sup>27</sup>. This ensures predictable performance and no interruptions (ideal for an enterprise use-case). You can choose CPU/RAM size for the VM based on app needs <sup>28</sup>. - **Scaling:** If the user base grows or workload increases, Replit's deployment system offers autoscaling (on certain plans). However, scaling might be limited compared to cloud platforms – an alternative is containerizing the app and deploying on a cloud if needed. But for initial implementation and moderate usage, a single Replit VM can suffice.



In summary, the architecture is a **typical web application**: a React front-end, an Express/Node back-end, and a PostgreSQL database – all of which can be developed and hosted on Replit. This separation of concerns will make the application maintainable and scalable, while Replit's cloud-based environment allows rapid iteration and testing.

## Rapid Development with Replit AI

One of Replit's strengths is its integrated AI tools (Replit Ghostwriter and Replit **Agent**) that can significantly accelerate development. We recommend leveraging these features to build the application faster:

- **Replit Agent for Project Setup:** Replit Agent can create full-stack apps from scratch by interpreting a natural language description <sup>29</sup>. You can start a new Repl using the Agent and provide a prompt like: *"Create a full-stack Node.js + React application for an Investment Approval Workflow. Include user authentication, a PostgreSQL database for storing users, roles, investments, requests, approvals, tasks, documents, and notifications. Implement two workflows (investment creation and cash request) with multi-level approvals and SLA tracking."* The Agent will scaffold the project structure in minutes – setting up the React app, the Express server, and integrating a database (it can even create initial schema).
- **Automated Database Integration:** Using the Agent, you can ask it to **add a PostgreSQL database** to the project if not already done. For example, *"Add a Postgres database with tables for users, roles, investment requests, cash requests, approvals, tasks, documents, notifications, and audit logs."* The Replit Agent can not only provision the database (hosted on Neon) but also generate the schema and necessary code to connect to it <sup>23</sup>. This can save a lot of time writing boilerplate model classes or migration scripts.
- **Using Ghostwriter for Code Generation:** Replit's Ghostwriter (AI code completion) can help write functions or components quickly. For instance, inside the Express app, you could write a comment or prompt like  

```
// Generate an Express route to create a new investment request with validation
```

and Ghostwriter may suggest code for you. While you'll need to review and test it, this speeds up mundane coding tasks.
- **Replit Import for UI Design:** If you have UI mockups (e.g., a Figma design of the screens), Replit's **Import** feature can convert designs into React code <sup>30</sup>. This could jumpstart your front-end development by producing the layout and styling as a starting point. You can then refine the components and add dynamic behavior.
- **Iterative Development via Chat:** The Replit Agent allows you to chat and make iterative requests. For example:
  - "Create a form component for New Investment with fields X, Y, Z and client-side validation."
  - "Implement JWT-based authentication for the Express app and protect all `/api` routes."
  - "Add a feature to send email notifications on new task assignment using NodeMailer."

The Agent will attempt to modify the codebase accordingly. Each request you make that results in code changes is saved as a **checkpoint** (with usage-based pricing if on a paid plan, but you only pay for

completed tasks) <sup>31</sup> <sup>32</sup> . You can roll back if something goes wrong, which makes experimentation low-risk.

- **Testing and Fixing with AI:** You can even ask the Agent or Ghostwriter to write tests (e.g., using Jest or a similar framework) for critical functions, or to fix a bug by describing the issue. For example, “The approval workflow isn’t moving to the next stage correctly when a manager approves – please fix this logic.” The AI might locate the issue and provide a solution, or at least guidance.
- **Documentation Generation:** Interesting tip – as seen in an anecdotal example, the agent can generate documentation for the code it wrote <sup>33</sup> <sup>34</sup> . You can prompt it to create a basic README or even API docs listing the endpoints, which is useful for handing over to stakeholders.

When using Replit’s AI, always review the generated code for correctness, security, and completeness. The AI can sometimes make assumptions or produce code that needs tweaking. However, it drastically reduces the grunt work. With complex requirements like ours, the AI might ask clarifying questions (e.g., “What fields should the Investment table have exactly?”) – providing it with the information (perhaps by copy-pasting relevant RFP text or your own detailed spec) will lead to better results. In essence, Replit AI tools allow the development team to focus more on **high-level design and validation** while the AI handles a lot of boilerplate coding.

## UI/UX Design – Screen-by-Screen Flow

A clear and intuitive user interface is crucial for adoption. Below is a proposed screen-by-screen flow, with key fields, validations, and navigation paths for each major screen in the application:

### 1. Login Screen

- **Purpose:** Authenticate users before they access the system.
- **Layout:** A simple login form with fields **Username** and **Password**, and a **Login** button. Possibly the company logo and system name at the top.
- **Fields & Validation:**
  - *Username* – could be email or a corporate username. Validate that it’s not empty and follows the expected format (e.g., a valid email format if using emails).
  - *Password* – validate not empty; for security, enforce a minimum length or complexity (though typically the backend handles password rules).
  - If credentials are wrong, show an error “Invalid username or password.”
- **Actions:** On clicking **Login**, the app calls `POST /api/auth/login` with the credentials. On success, the user’s session/JWT is stored, and they are navigated to the **Dashboard**. On failure, an error message is displayed.
- **Navigation:** Also include options like “Forgot Password?” (which might redirect to a recovery flow, if implemented) or information on how to get an account (if accounts are provisioned by admins, you might not have self-signup).

### 2. Dashboard / Home Page

- **Purpose:** Provide a summary and quick access to key sections for the logged-in user.
- **Content:** The Dashboard can show:

- **My Pending Tasks:** a widget or table listing the top 5 tasks waiting for the user's action (with a link to view all tasks). Columns: Task name/description (e.g., "Approve Investment X"), due date (if SLA), and an "Open" action.
- **My Requests Status:** a widget showing the user's recently submitted requests and their statuses (e.g., "Investment ABC – Pending Manager Approval").
- **SLA Alerts:** if the user is a manager/admin, maybe show any requests in their team that are overdue.
- **Navigation Cards/Links:** Big buttons or menu links to "New Investment Request", "New Cash Request", "View All Investments", "My Tasks", etc. This can be in a sidebar or top menu as well.
- **Interactions:** Clicking on a task will navigate to the **Approval/Detail screen** for that request. Clicking on a request status could open the detailed view of that request.
- **Role-Based Content:** The Dashboard can be customized. For example, an Admin might see system stats or admin actions (user management links) instead of personal requests. A Finance user might see pending cash disbursement tasks.

### 3. New Investment Request Form

- **Purpose:** Allow users to submit a new investment opportunity for approval.
- **Fields:** (These reflect the required data mentioned in RFP <sup>3</sup>)
- **Investment Title/Name** – Short text identifying the opportunity.
- **Investment Company** – The company or issuer of the investment (e.g., if investing in a fund or company, specify the name). Could be a dropdown if limited to known companies or free text.
- **Investment Type** – Dropdown or options (e.g., Equity, Fixed Income, Real Estate, etc., based on the types defined by Tawuniya).
- **Amount / Value** – Numeric field for the investment amount (validate for positive number, possibly currency format).
- **Risk Level** – Perhaps a dropdown (Low, Medium, High) or a numeric risk score. Could be filled by the analyst based on internal risk assessment.
- **Description / Rationale** – Text area for the investment thesis or description.
- **Attachments** – File upload for documents (at least the proposal document). Allow multiple files. Validate file types (pdf, docx, xlsx perhaps) and size (maybe up to a few MB each).
- **Other Fields:** Depending on needs, could include: expected ROI, investment horizon (duration), any compliance checklist, etc. (The RFP mentions "all required data"; we list the obvious ones, but an actual form might be more extensive.)
- **Validations:** All key fields (title, type, amount, company) should be required. Amount should be numeric and within reasonable bounds (and maybe currency selection if multi-currency). Attachments could enforce at least one document attached (if that's a requirement to proceed with approval).
- **Layout:** Use a multi-section form if it's long. Possibly separate into sections: Basic Info, Financial Info, Risk & Attachments.
- **Submission:** The form has **Submit** and maybe **Save as Draft**. Draft functionality would allow saving an incomplete form to finish later (optional feature). On submit, do client-side validation then call `POST /api/investments` with the data. Show a confirmation or redirect to the newly created request's detail page.
- **Post-Submission:** After successful creation:
  - The user is redirected to an **Investment Detail page** for that request, showing it in "Pending" status and waiting for approvals.
  - A success message "Investment request created successfully" is shown.

- The workflow engine on backend will have created tasks for approvers; the UI might proactively inform the user like “Waiting for Manager Approval”.

#### 4. Investment Request Detail / Approval Screen

- **Purpose:** Display a specific investment request in detail – used by both requestors (to check status) and approvers (to review and take action).
- **Content:**
  - **Summary Header:** Top section showing key info like Request ID, Title, current Status (“Pending Manager Approval” or “Approved” etc.), and maybe a badge or progress indicator of what stage it’s in (e.g., 1 of 3 approvals completed).
  - **Details Sections:** All the fields submitted (company, type, amount, description, etc.) displayed in read-only form.
  - **Attachments:** List of documents attached. Each with a download link. Possibly allow preview if it’s an image/PDF (for convenience, an embedded viewer).
  - **History/Comments:** A timeline of actions – e.g., “Created by Alice on Jan 10”, “Approved by Manager Bob on Jan 11 (comment: looks good)”, etc. This audit trail is important for transparency.
  - **Approval Actions (for approvers):** If the current user has an outstanding task to approve this request, they should see **Approve** and **Reject** buttons (or “*Send Back for Revision*” if that’s a path). They should also have a text box to add comments when taking an action. (Comments are especially important on rejection or send-back, to explain why).
  - **Edit (for requestor, if allowed):** If the request is in a state where the requestor is allowed to make changes (e.g., it was sent back or maybe still draft), show an **Edit** button. This opens an edit form (similar to the create form) pre-filled with data.
- **Validations/Behaviors:**
  - Approvers must enter a comment if rejecting (enforce via UI).
  - Once an approver takes action, disable the buttons to prevent duplicate actions.
  - If the request is approved and fully completed, no action buttons are shown; just display “Approved on [date]”.
- **Navigation:**
  - Approvers might arrive here by clicking a task from their list. After taking action, they could be redirected back to their task list (with a success message), or stay on the page which now shows the updated status.
  - The requestor might navigate here from “My Requests” list to see the status. They might also get here right after submitting.
  - Provide a **Back** link to go back to the previous list (either tasks or requests list).
- **Notifications:** Optionally, from this page an approver might have a “Remind” button if they need to ping someone (but usually automated). Not a requirement, but sometimes present in workflow apps.

#### 5. New Cash Request Form

- **Purpose:** Allow users to request cash for an existing investment.
- **Access:** This form might be accessed either from a specific Investment’s page (e.g., “Request Cash” button on an approved investment) or from a general “New Cash Request” menu which then asks the user to select the investment.
- **Fields:**

- **Investment Selection:** If coming from an Investment page, this can be pre-filled/hidden. Otherwise, a dropdown or autocomplete field to select which investment the cash is for (limited to investments the user has access to, likely their own or their team's, or all if finance).
- **Amount:** The cash amount requested (numeric, required).
- **Purpose/Description:** Text explaining why the cash is needed (e.g., "Funding Q1 phase of project").
- **Attachments:** If any supporting docs (maybe an invoice or payment schedule).
- **Date Needed By:** Possibly a date field if they need the funds by a certain date (to inform SLA for finance).
- **Validation:** Amount required and >0. If the amount is more than the remaining budget of that investment (if such data exists), could warn or disallow (business rule). Purpose required to avoid blank requests. Date needed – should be today or future.
- **Submission:** On submit, calls `POST /api/cashrequests`. After creation:
  - Manager approver receives a task to approve this cash request.
  - The requestor is redirected to a **Cash Request Detail** (similar to Investment detail but simpler).
- **Cash Request Detail:** Would include fields similar to the investment detail but fewer: shows the cash request info and its status, plus any approvals. Likely a one-step approval (manager -> then notify finance).
- If approved, also display that it was forwarded to Finance, with maybe a field like "Finance Status" (which an accountant can mark as "Paid" once done).
- Approver sees Approve/Reject buttons on this detail if pending.

## 6. Task List (My Tasks)

- **Purpose:** Provide users (especially approvers) a consolidated view of all tasks assigned to them.
- **Layout:** Likely a table with columns: **Task** (a descriptive name like "Approve Investment – Project X" or "Review Cash Request – Investment Y"), **Type** (Investment or Cash), **Requested By** (who initiated, for context), **Requested Date** (or due date), **SLA Remaining** (could show time left or overdue status), and an **Action** link/button ("Open" to go to detail).
- **Filtering/Sorting:** Users should be able to filter tasks by type or due date, and sort e.g. by oldest first or by investment name. This is useful if there are many tasks.
- **Batch Actions:** Possibly allow multi-select and bulk actions if, say, an approver wants to approve multiple in one go (though in our scenario, likely each requires individual attention, so maybe skip bulk actions).
- **Navigations:** Clicking open goes to the detail screen. We should allow easy navigation back to the task list from the detail (a back link, or the list could open detail in a modal – but that can be complex; simpler to navigate).
- **Visual Cues:** Overdue tasks could be highlighted in red or flagged. Tasks nearing SLA could have a warning icon.

## 7. My Requests / Investment List

- **Purpose:** Let users see all requests they have submitted (both investments and cash requests, or separate lists for each).
- **Layout:** Possibly two tabs: "My Investment Requests" and "My Cash Requests". Or a single table with a Type column.
- **Columns:** **Request** (maybe a name or ID), **Type**, **Status** (Pending, Approved, Rejected, etc.), **Last Updated** or Current Stage ("With Investment Committee" for example), **Actions** (View detail, and if status is Draft or Returned, maybe Edit).

- **Search/Filter:** Allow searching by name or filtering by status (so a user can find old approved investments or see all rejected ones).
- **Navigation:** Clicking on a request opens the detail page. The user can also click “New Investment” or “New Cash” from here if they want to create another.

## 8. Administration Screens (for Admin Role)

If an administrator role is present, an **Admin Panel** or section should be available. Key admin screens:

- **User Management:** A screen listing all users, with the ability to add new users, assign roles, deactivate users, etc. Fields for user creation: Name, Email/Username, Role(s), etc. For existing users, an edit form to change roles or reset password. (If integrating with SSO, user management might be limited to viewing, as users sync from an identity provider.)
- **Role Management:** Define roles and their permissions. This could be as simple as a predefined list in code (not all apps expose UI for this). But at least a reference of what each role can do, maybe in documentation.
- **Workflow Configuration:** If allowing dynamic workflows, an interface to configure the approval steps:
  - For example, an “Investment Workflow” editor where an admin can list out steps: (Level 1: Role = Investment Manager, SLA = 2 days; Level 2: Role = Investment Committee, SLA = 3 days; etc.). The admin can add/remove levels or change SLA times. This could be stored in a config table.
  - Similarly for “Cash Request Workflow”.
  - If not building a full UI for this, at least store in DB and maybe allow tech staff to update. But a visual editor is a nice-to-have (though complex to build fully).
- **Audit Logs:** A viewer for audit trails (see Audit Trails under Best Practices). The admin can search logs of who did what, when. This is useful for compliance and debugging issues (e.g., “Who approved this request? Who changed this field?”).
- **System Settings:** Possibly a page for other configurations – e.g. define SLA durations (if not fixed per workflow), toggle maintenance mode, etc. This depends on how many dynamic settings we allow.

**Note:** For the first iteration, not all admin features need a UI (some settings can be in database or config). Focus might be on user management since that’s critical to manage access.

## Screen Flow Summary Table

For clarity, the following table summarizes the main screens, their purpose, and key elements:

Screen	Purpose	Key Fields/Elements	Actions/Navigations
<b>Login</b>	Authenticate users	Username, Password fields; Login button	On success -> Dashboard; error messages
<b>Dashboard</b>	User’s overview & navigation	Pending Tasks widget; Recent Requests; Navigation links	Links to Task list, New Request forms, etc.

Screen	Purpose	Key Fields/Elements	Actions/Navigations
<b>New Investment Request</b>	Submit new investment for approval	Form fields: Title, Company, Type, Amount, Risk, Attachments...	Submit -> creates request -> Investment Detail; validations on fields
<b>Investment Detail</b>	Show request info; approvers act here	Request details, attachments, history, approval buttons	Approvers: Approve/Reject; Requestor: possibly Edit if allowed; Back to list
<b>New Cash Request</b>	Submit cash drawdown for an investment	Form fields: select Investment, Amount, Purpose, Attachments	Submit -> creates cash request -> Cash Detail
<b>Cash Request Detail</b>	Show cash request status & approval	Cash request details, status, approvals (manager & finance)	Manager: Approve/Reject; Finance: mark paid (if applicable)
<b>My Tasks</b>	List tasks waiting for user action	Table of tasks (name, type, due, from who)	Open task -> goes to Detail; filter/sort tasks
<b>My Requests</b>	List user's submitted requests (all types)	Table of requests (ID, type, status, stage, last update)	View -> goes to Detail; maybe Edit for drafts
<b>Admin: Users</b>	Manage user accounts and roles	List of users; form to add user; edit role assignments	Save new users; change roles; reset password
<b>Admin: Workflows</b>	Configure approval workflow steps (optional)	For each process, define approval levels and SLA	Add/remove levels; Save config to DB
<b>Admin: Audit Logs</b>	View change logs for compliance	Table of audit entries (timestamp, user, action, details)	Filter by date/user; export logs if needed

(The above is a conceptual summary; the actual UI can be refined based on user feedback and usability testing.)

## Database Schema Design

A robust **PostgreSQL schema** underpins the application, capturing all necessary entities and their relationships. Below we propose tables for each major entity listed (users, roles, investments, workflows, approvals, tasks, documents, notifications, audit trails), along with key fields. (For brevity, only main fields are shown, omitting trivial fields like foreign key IDs in some descriptions where context implies.)

### Users and Roles

We'll use a standard user-role structure to support permissions:

Table: users

Field	Type	Description
id	serial (PK)	Unique user ID.
username	text	Login username (or email). Must be unique.
password_hash	text	Hashed password (if not using external auth).
name	text	Full name of the user.
email	text	Email address (could double as username).
active	boolean	Whether the account is active (for disabling).
created_at	timestamp	When the user account was created.
updated_at	timestamp	Last update timestamp.

**Table:** roles

Field	Type	Description
id	serial (PK)	Unique role ID.
name	text	Role name (e.g., "Analyst", "Manager", "Admin").
description	text	Description of the role's purpose.

**Table:** user\_roles

Field	Type	Description
user_id	integer (FK -> users.id)	Reference to a user.
role_id	integer (FK -> roles.id)	Reference to a role.
PRIMARY KEY	(user_id, role_id)	Composite PK (one user can have multiple roles).

*Note:* If each user has exactly one role, we could instead put a role\_id in `users` table. But using a join table `user_roles` offers flexibility for multi-role assignments. Roles will define what a user can do in the application logic (not enforced by the DB itself beyond just grouping roles).

## Investment Requests

This table holds the main data for investment opportunity requests (both pending and approved investments):

**Table:** investments (could also be named `investment_requests` for clarity)



Field	Type	Description
id	serial (PK)	Unique ID for the investment request.
title	text	Title or short name of the investment.
company	text	Name of the company or entity of the investment.
investment_type	text	Type/category of investment (could be FK to a types lookup table).
amount	numeric(15,2)	Proposed investment amount (with two decimal places for currency).
currency	text	Currency code (e.g., "USD", "SAR") if multi-currency support needed.
risk_level	text or smallint	Risk rating/level (could be "Low/Med/High" or a numeric score).
description	text	Detailed description or rationale.
status	text	Current status of request (e.g., "Draft", "Pending Approval", "Approved", "Rejected").
created_by	integer (FK -> users.id)	User who created the request.
created_at	timestamp	Submission time of the request.
updated_at	timestamp	Last updated time (e.g., if edited or status change).
approved_at	timestamp (Nullable)	Time the request was fully approved (if it is).
approved_by	integer (FK -> users.id, Nullable)	The final approver (if applicable).

Additionally, there may be fields like `rejected_at`, `rejected_by`, or a separate table for approval history (covered in Approvals table) instead of storing only final approval. The `status` will reflect if pending and how many stages done; it could be granular (e.g., "Pending Level 2") or just high level.

## Cash Requests

We keep a separate table for cash requests since they are a different workflow but related to investments:

**Table:** `cash_requests`

Field	Type	Description
id	serial (PK)	Unique ID for the cash request.
investment_id	integer (FK -> investments.id)	The related investment this cash is for.

Field	Type	Description
amount	numeric(15,2)	Cash amount requested.
currency	text	Currency (should likely match the investment's currency).
purpose	text	Description/purpose of the cash usage.
status	text	Current status (e.g., "Pending Approval", "Approved", "Rejected", "Disbursed").
created_by	integer (FK -> users.id)	Who created the cash request (likely the investment owner).
created_at	timestamp	Request submission time.
updated_at	timestamp	Last update time.
approved_at	timestamp (Nullable)	When manager approved (if approved).
approved_by	integer (FK -> users.id, Nullable)	Manager who approved.
finance_notified_at	timestamp (Nullable)	When finance was notified (likely same as approved_at).
finance_processed_at	timestamp (Nullable)	When finance marked as processed (if we track that).
finance_processed_by	integer (FK -> users.id, Nullable)	Finance user who marked processed.

The `status` might go through "Pending Manager Approval" -> "Approved/Pending Finance" -> "Completed". We include fields to track approvals and processing times.

## Workflow and Task Management

To manage approvals and tasks, we design two tables: one for the definition of **approval steps** (if needed), and one for the **task instances** assigned to users. Additionally, an **approvals log** table will capture each approval decision (this could be combined with tasks, but we separate for clarity).

**Table:** `workflow_steps` (optional; defines template for workflows)

This table defines the sequence of approvals for a given workflow type. For example, an Investment workflow might have 3 steps, Cash workflow 1 step. This allows dynamic reconfiguration.

Field	Type	Description
id	serial (PK)	Unique step ID.
workflow_type	text	Type of workflow (e.g., "Investment", "Cash").

Field	Type	Description
step_number	integer	The order of this step (1, 2, 3...).
role_id	integer (FK -> roles.id)	Role responsible for this step (e.g., Investment Manager role).
step_name	text	Optional name (e.g., "Manager Approval", "Committee Approval").
sla_days	integer (Nullable)	SLA in days for this step (or hours, depending on granularity).
notify_role	integer (FK -> roles.id, Nullable)	Role to notify upon completion (e.g., notify Finance at end).

For cash requests, there might be one step (manager) and notify Finance. For investments, multiple steps. If the workflow is static and coded, this table might not be necessary; but it's useful for admin configurability.

**Table:** `tasks` (approval tasks assigned to users)

Field	Type	Description
id	serial (PK)	Unique task ID.
task_type	text	Type of task ("InvestmentApproval" or "CashApproval", or perhaps "GeneralTask").
investment_id	int (FK -> investments.id, Nullable)	Linked investment (if task is for an investment approval).
cash_request_id	int (FK -> cash_requests.id, Nullable)	Linked cash request (if task is for cash approval).
assigned_to	int (FK -> users.id)	User who must do the task (could be null if unassigned and claimable by role).
assigned_role	int (FK -> roles.id)	Role expected to handle this task (useful if not yet assigned to specific user).
step_number	int	The workflow step this corresponds to (e.g., 1 for first approval).
status	text	"Pending", "Completed", "Withdrawn", etc.
due_date	timestamp (Nullable)	Deadline for SLA (calculated from created_at + sla).
created_at	timestamp	When task was created/assigned.
completed_at	timestamp (Nullable)	When task was completed (if it is).

Field	Type	Description
completed_by	int (FK -> users.id, Nullable)	Who completed it (in case tasks can be completed by someone else, e.g., an admin on behalf).
outcome	text (Nullable)	Result of task, e.g., "Approved", "Rejected", or other resolution.
comments	text (Nullable)	Comments provided upon completion (like approver's notes).

This `tasks` table essentially tracks the *work items*. For an investment with 3 approval steps, it will have up to 3 task records (one for each approver in sequence, created as the process moves along). For a cash request, maybe 1 task record. If tasks are used for things beyond approvals (like "provide more info" tasks), those could also be represented here by a different `task_type`.

**Table:** `approvals` (log of approval decisions)

This table logs each approval/rejection action taken. (One could consider the `tasks` table's completed records as implicitly the approvals log, but we can have a separate table for clarity and historical record even for rejected flows.)

Field	Type	Description
id	serial (PK)	Unique approval action ID.
investment_id	int (FK -> investments.id, Nullable)	Related investment (if this approval was for an investment).
cash_request_id	int (FK -> cash_requests.id, Nullable)	Related cash request (if for cash).
approver_id	int (FK -> users.id)	User who made the approval decision.
role_id	int (FK -> roles.id)	Role of the approver (for context, e.g., was acting as Manager or Committee).
step_number	int	Which step this was in the workflow (if applicable).
decision	text	"Approved" or "Rejected". (Could also include "Revised" etc.)
comments	text	The comment the approver left.
decided_at	timestamp	Timestamp of the decision.

By logging approvals separately, we keep a permanent record even if a request is later edited or resubmitted. This is helpful for audit trails (who approved what and when).

## Documents

Documents (attachments) will be stored in Object Storage, but we maintain a table to track them:

**Table:** documents

Field	Type	Description
id	serial (PK)	Document record ID.
investment_id	int (FK -> investments.id, Nullable)	If the document is related to an investment request.
cash_request_id	int (FK -> cash_requests.id, Nullable)	If related to a cash request.
file_name	text	Original file name (for display).
file_type	text	MIME type (e.g., application/pdf).
file_size	integer	Size in bytes.
storage_url	text	URL or key for file in object storage (could be a Neon bucket ID or GCS link).
uploaded_by	int (FK -> users.id)	Who uploaded the file.
uploaded_at	timestamp	Upload timestamp.
description	text (Nullable)	Optional description or label for the document.

We include foreign keys for both possible parent entities (investment or cash request). Only one of those would be non-null for a given record. Alternatively, we could have a generic `request_type` and `request_id` to link to any type of request if we had more types.

Using Replit's Object Storage via bucket, each file could be stored with a key structure like: `investment-{id}/{filename}` or similar for organization. The `storage_url` might be a signed URL generated on the fly rather than stored, but a base path or key is stored.

## Notifications

For tracking notifications (especially for audit and maybe for in-app display):

**Table:** notifications

Field	Type	Description
id	serial (PK)	Notification ID.

Field	Type	Description
user_id	int (FK -> users.id)	The user who the notification is for.
message	text	The content of the notification (brief text).
link	text (Nullable)	Optional link (URL or route) that the notification relates to (e.g., link to the request detail).
is_read	boolean	Has the user marked/viewed it as read.
created_at	timestamp	When the notification was created/sent.
type	text	Type of notification (e.g., "TaskAssigned", "Reminder", "StatusUpdate").
related_entity	text (Nullable)	Type of entity related (if needed to decode link, e.g., "Investment").
related_id	int (Nullable)	ID of the related entity (e.g., Investment ID).

This table allows showing a notification list in the UI, and also records that say an email was sent (if we log that as well). In-app, once the user views it, the front-end would mark it as read.

## Audit Trails

To comply with corporate governance and to ease debugging, an **audit trail** of changes and actions is essential:

**Table:** `audit_logs`

Field	Type	Description
id	serial (PK)	Audit log entry ID.
timestamp	timestamp	Time of the action.
user_id	int (FK -> users.id, Nullable)	User who performed the action (nullable if system).
entity_type	text	The type of entity changed (e.g., "Investment", "CashRequest", "User").
entity_id	int (Nullable)	The ID of the entity affected (if applicable).
action	text	Description of action (e.g., "STATUS_UPDATE", "CREATE", "LOGIN", "EDIT_FIELD").
details	text (Nullable)	Additional info (could be JSON or text describing the change, e.g., "status from X to Y", or the changed fields and values).

This table can grow large, so consider retention policies or archiving after some time. But it is invaluable for tracing what happened in the system. E.g., if an investment's amount is edited or a task reassigned, log it here.

## Schema Implementation Notes

- **Foreign Key Constraints:** Ensure to add foreign key constraints for referential integrity (with ON DELETE rules as appropriate). For example, if an investment is deleted (though typically you wouldn't truly delete – probably just mark inactive), you might want cascading deletes on related tasks, documents, etc., or prevent deletion if related records exist.
- **Indexes:** Index common query fields:
  - e.g., `tasks(assigned_to, status)` for quickly finding pending tasks for a user.
  - `investments(created_by)` and `cash_requests(created_by)` for listing a user's requests.
  - `audit_logs(entity_type, entity_id)` for retrieving logs per entity.
  - Full-text index on investment title or description if searching is needed.
- **Unique Constraints:** e.g., user username or email unique, perhaps a constraint that one investment can't have two active cash requests of same exact parameters (business-specific).
- **Partitioning:** Not needed initially, but if audit\_logs or notifications get huge, could consider partitioning by date for performance.

## Compliance, Security, and Performance Best Practices

Building an enterprise application requires attention to non-functional requirements as well. Below are best practices in the context of Replit and our application:

### Security & Data Protection

- **Authentication & Session Security:** Implement secure authentication (preferably JWT for SPA or secure HttpOnly cookies for sessions). Use strong password hashing (if managing passwords) and consider **2FA** if needed for sensitive approvals. Replit provides a secret storage for environment variables (like DB passwords, API keys) – **never hard-code secrets** in the code <sup>35</sup>. All communication should be over HTTPS (Replit by default serves over HTTPS).
- **Authorization & RBAC:** Enforce role-based access on the backend for every endpoint. For example, only admins can call admin endpoints; only an approver can call the “approve request” API for a request assigned to them, etc. Do this server-side (don't rely solely on hiding buttons in the UI).
- **Input Validation & Sanitization:** All inputs from forms or API should be validated on both frontend (for user experience) and backend (for security). This prevents bad data or malicious input. Utilize frameworks or libraries to sanitize inputs against SQL injection and XSS. If using an ORM, many do parameter binding which mitigates SQL injection. Also sanitize output displayed on pages to avoid reflecting any injected script (if using React, it escapes content by default, unless using `dangerouslySetInnerHTML` which we should avoid or use carefully).
- **CSRF Protection:** If using cookie-based auth, implement CSRF tokens for state-changing requests. If using JWT in an SPA, CSRF is less of an issue (as tokens are in JS memory), but then protect against XSS which could steal tokens. Consider using HttpOnly cookies for JWTs if possible to mitigate that (and then need CSRF tokens).
- **Secure File Handling:** When users upload files, store them in a secure location (Replit's Object Storage uses Google Cloud Storage with robust security <sup>16</sup>). Ensure files are scanned for malware if

needed and that only authorized users can access them (no open bucket access – use signed URLs that expire, or authenticated download endpoints). Also validate file type by checking content (to prevent someone uploading an executable and then somehow executing it).

- **Encryption:** While our database (Neon Postgres) is managed by Replit/GCP and data is encrypted at rest by the provider <sup>36</sup>, we should ensure sensitive data in our domain (like maybe investment documents or personal info) is encrypted if needed at application level. Likely not necessary for this use-case beyond the standard encryption in transit (TLS) and at rest provided by platform. All Replit hosted traffic is HTTPS <sup>37</sup>, and DB connections should use SSL as well.
- **Audit and Monitoring:** Use the audit\_logs to periodically review if any unusual activity (e.g., unauthorized access attempts, or approvals done out of order). Consider integrating logging to an external SIEM if this goes to production (optional at MVP stage).
- **Secrets Management:** Use Replit's Secrets for storing DB credentials, API keys (e.g., email service API). Replit's environment ensures these are not exposed and are injected as env vars <sup>24</sup>. Rotate credentials if needed and handle them carefully in code (avoid printing in logs).
- **Platform Security:** Replit itself runs on GCP and has obtained **SOC 2 Type II** compliance <sup>38</sup>, and complies with GDPR/CCPA for data privacy. This provides a baseline of trust. They also have internal security teams and WAF protections at the platform level <sup>39</sup> <sup>35</sup>. However, as a developer, we must still build the app securely as above.

## Compliance and Regulatory Considerations

- **Audit Trails:** As described, maintain detailed audit logs. This helps with internal audits and any regulatory compliance (e.g., if this system falls under financial controls or needs to demonstrate who approved what).
- **Data Retention Policies:** Determine how long data should be kept. Investment records likely need long-term retention. Logs might be pruned after X years if not needed. Ensure data disposal (if needed) is done properly.
- **Privacy:** Although this system mostly handles internal corporate data (not personal customer data), if any personal data is stored (even employee names, etc.), comply with relevant privacy regulations. Likely minimal concerns here aside from securing data.
- **Compliance Checks:** The RFP's broader scope (outside item 1) mentions compliance and regulatory reporting. If the system integrates those, ensure any regulatory reports (like to regulators) have accurate data. Possibly out of scope for item 1 implementation guide.
- **Segregation of Duties:** In workflow design, ensure that one user cannot both initiate and approve an investment in a conflict of interest manner. The roles structure should enforce that (e.g., an analyst cannot approve their own request; the system should prevent it by not assigning them as approver and by blocking if they try any API misuse).
- **E-Signature or Approvals Authenticity:** If the company requires, you might need to ensure that an approval action is tied to the user (which we do by login sessions). For very high assurance, two-factor auth or signing might be needed for final approvals, but that's beyond normal scope. Logging user ID and requiring secure login is typically sufficient.

## Performance & Optimization

- **Efficient Queries:** As usage grows, avoid N+1 query problems. Use joins or batch queries to load data. For example, when showing a list of requests with their current approver, join the tasks or user tables instead of per-row lookups. Optimize with indices as noted.



- **Caching:** Use caching cautiously – e.g., if there is data that doesn't change often (like a list of investment types or user info), the backend can cache it in memory to avoid frequent DB hits. Could also use a Redis cache (Replit might not directly provide Redis, but perhaps use the key-value store or an in-memory JS cache).
- **Pagination:** Implement pagination or lazy loading on tables (tasks list, requests list) to avoid heavy loads if there are thousands of records.
- **SLA and Background Jobs:** If SLA monitoring needs to send a notification exactly when overdue, you might need a background scheduler (to check and send alerts). On Replit, you could either:
  - Use a cron job (Replit doesn't run cron when asleep unless Always On; with Always On or Reserved VM it could).
  - Or every time a list loads, update statuses (not ideal). Possibly use the database's ability to schedule a function or simply rely on frequent checks.
  - Simpler: each time a page is loaded, if something is overdue by comparing timestamps, trigger an email if not already sent. Or have a small setInterval in the backend.
  - For robust solution, integrate a job scheduler (there are Node libraries like node-cron or Agenda).
- **Concurrent Usage:** The app should handle multiple simultaneous users. Node/Express can handle many concurrent requests, but ensure no shared global state is modified unsafely. The PostgreSQL can handle concurrent transactions; just be mindful of transaction isolation if needed (like two people approving the same thing at once – one should ideally be prevented because once one approves, the task is completed).
- **File Upload Performance:** For large files, consider direct upload to storage (S3/GCS) from the client to offload the server. But Replit's object storage likely requires going through the app. Perhaps fine if file sizes are moderate. Ensure to stream files to storage rather than reading whole into memory, to keep memory usage low.
- **Testing at Scale:** Use test data to simulate, say, hundreds of requests and tasks to see if any page is slow. Optimize queries or add indexes as needed.

## Deployment Readiness and Limitations on Replit

Finally, understand the **deployment environment** limitations and prepare the app accordingly:

- **Always On vs. Reserved VM:** As mentioned, use Always On or a Reserved VM deployment to ensure the application doesn't sleep <sup>27</sup>. A sleeping app would cause users to experience delays or missed notifications.
- **Resource Limits:** Replit (depending on plan) has limits on CPU, memory, and storage. Enterprise plans allow higher limits and even private clusters. But a typical Repl might have a few GB of storage and some RAM. Monitor resource usage (Replit provides metrics) and consider upgrading or optimizing if hitting limits. For example, if heavy PDF documents are stored, ensure the bucket usage is accounted for (object storage costs by usage <sup>40</sup> <sup>41</sup>).
- **Scalability:** Replit's infrastructure can scale your app to an extent (especially if using Deployments with autoscale on their hosting). However, it's not the same as having a cluster of servers behind a load balancer that you might get on AWS/Azure. For moderate enterprise use (let's say tens of concurrent users, hundreds of requests/day), Replit should handle it. If you anticipate very high load, you might need to plan for horizontal scaling (which may mean moving to a more traditional cloud setup or waiting for Replit's future enterprise features).
- **Vendor Lock & Migration:** Building on Replit is fast and convenient. But note that the code is portable – you can always take the Node/React app and deploy elsewhere if needed (since it's standard technologies). Keep the architecture fairly provider-agnostic (don't use too many Replit-

exclusive APIs beyond convenience tools). For example, using Neon Postgres is fine since you can also host Postgres elsewhere if migrating. Using Replit's object storage (GCS) is okay, just note data export if moving.

- **Continuous Deployment:** Replit can connect to GitHub or support Git, but typically the Repl itself is your environment. For enterprise robustness, you might use Replit Teams for version control or export code to a repository for backup. Test thoroughly in the Replit workspace, then use the **Deployments** feature to push the stable version to production (which then runs isolated from your dev workspace).
- **Custom Domain & SSL:** If this needs to be on a company domain (e.g., `ims.tawuniya.com`), Replit allows adding custom domains to deployments <sup>28</sup> <sup>42</sup>. Ensure to set that up so that it's accessible at a professional URL with SSL (Replit provides SSL certificates automatically for custom domains).
- **Data Persistence:** Reiterate that any data that needs to persist must be in the PostgreSQL database or object storage. The Repl's local filesystem (other than the `/mnt/data` persistent mount, which might be limited) is not meant for long-term storage of app data <sup>43</sup>. So, for example, do not just save uploaded files to disk – they'd be gone if the container restarts or shifts. Use the database and cloud storage as designed.
- **Third-Party Services:** Check that any integration (like email sending or SSO) can work from within Replit (outbound internet access is allowed in Repls, so hitting an API or SMTP server is fine). If company firewalls are a consideration (since Tawuniya might prefer the app running in a controlled network), Replit might need to be vetted by their IT. Replit being cloud-based means you'll want to ensure it meets any internal security criteria.
- **Testing & QA:** Before going live, do thorough testing. Replit makes it easy to fork the Repl or create a separate staging deployment. Use test accounts for each role to simulate the workflows. Ensure that the workflows function as intended (all the way from submission to multi-level approvals and closure). Also test failure cases (rejections, SLA triggers, etc.).

In conclusion, **Replit provides an excellent rapid development environment** to build this investment approval workflow application from scratch. By following the above architecture and best practices, we can meet the functional requirements (investment creation, cash requests, SLA tracking, document management, tiered approvals <sup>3</sup> <sup>4</sup>) while also ensuring the application is secure, compliant, and performs well. With a carefully designed UI and robust backend, the system will streamline Tawuniya's investment processes and improve visibility and control, aligning perfectly with the goals stated in the RFP <sup>44</sup> <sup>45</sup>.

#### Sources:

- Tawuniya RFP – Investment Management System, Item 1 (Investment Operations & Workflows) <sup>3</sup> <sup>4</sup> <sup>2</sup>
- Replit Documentation – Database (PostgreSQL via Neon) and AI Agent usage <sup>23</sup> <sup>29</sup>
- Replit Documentation – Object Storage for file uploads (uses GCS) <sup>16</sup>
- Replit Platform Security & Compliance (GCP hosting, SOC2, secure development practices) <sup>39</sup> <sup>35</sup>
- Replit Deployment Options – Always On and Reserved VM for production stability <sup>46</sup> <sup>27</sup>

16 Replit Docs

<https://docs.replit.com/cloud-services/storage-and-databases/object-storage>

19 35 36 38 39 40 41 Replit: An Analysis of the AI-Powered Cloud Development Platform

<https://www.baytechconsulting.com/blog/replit-an-analysis-of-the-ai-powered-cloud-development-platform>

20 Ansh & Riley's Template 3 (Full-stack React & Node.js) - Replit

<https://replit.com/@an732001/Ansh-and-Rileys-Template-3-Full-stack-React-and-Nodejs>

21 22 23 24 Replit Docs

<https://docs.replit.com/cloud-services/storage-and-databases/sql-database>

25 26 30 37 43 46 Replit — Hosting Apps with Always On

<https://blog.replit.com/alwayson>

27 28 42 Replit Docs

<https://docs.replit.com/cloud-services/deployments/reserved-vm-deployments>

29 31 32 Replit Docs

<https://docs.replit.com/replitai/agent>

33 34 Down the Rabbit Hole with Replit Agent: Building a Full-Stack Application Without Writing a Single Line of Code

<https://www.linkedin.com/pulse/down-rabbit-hole-replit-agent-building-full-stack-single-lahullier-llvje>